

Predicting Adverse Thermal Events in a Smart Building*

Jenna MacCarley
Carnegie Mellon University
jmaccarl@andrew.cmu.edu

Federico Ponte
Carnegie Mellon University
federico.ponte@sv.cmu.edu

Victor Hu
Carnegie Mellon University
vyh@andrew.cmu.edu

Diane Yang
Carnegie Mellon University
dan.yang@sv.cmu.edu

ABSTRACT

The purpose of this research is to use machine learning methods to create an effective anomaly detection system, designed to predict real adverse thermal events in a smart building located on NASA’s Sustainability Base. Our project centers around the design of an effective anomaly detection system modeled after ACCEPT, a state of the art system built by NASA engineers for comparing various regression models and detection methods. Our goal was to create a similarly performing system, but use open-source Python libraries as opposed to proprietary software based on MATLAB. Ultimately, we were able to create a highly effective system that could predict a cold event’s occurrence using the given data with a 0.6% False Positive rate and a 3.1% Missed Detection rate. We compared our research results to those from another student team from Cornell, that used ACCEPT to predict cold events with the same data, and found that our system design results rivaled theirs. The source code for our research is publicly available online¹.

1. INTRODUCTION (*JENNA*)

NASA’s Sustainability Base [28] is a green building located in the Ames Research Center at Moffet Field, CA and among the greenest in the federal government. It is LEED platinum certified, achieving the highest rating for sustainability through its innovative design of architecture, resource recycling systems, renewable power generation and advanced sensing systems. The advanced sensing systems in the building are advertised as “anticipating and reacting to changes in sunlight, temperature, wind, and occupancy and will be able to optimize its performance automatically, in real time, in response to internal and external changes” [28]. However, despite these advancements in technology there have been numerous complaints that rooms in the building are too cold. Since the building incorporates thousands of sensors, including 2636 thermal sensors specifically [4], it is difficult

to narrow down where this issue of cold complaints may be coming from.

Our goal for this research is to use machine learning models to understand what events may be triggering an abnormally cold temperature in the building and be able to predict instances of it. Our research is based off of a project out of Cornell university, which had the same goal of predicting adverse thermal events in this building [4]. These researchers used ACCEPT (Adverse Condition and Critical Event Prediction Toolbox), a MATLAB framework developed by NASA which is designed to compare and contrast the performance of various machine learning algorithms and anomaly detection mechanisms [18]. Our goal for this project is to expand upon their research, but develop our own framework in Python using existing machine learning libraries such as scikit-learn [27] and Pyflux [23]. Our goal was to ultimately replicate and improve upon the results from the Cornell researchers, creating higher accuracy in the prediction of the anomalous cold events determined by the resulting false positive and missed detection rates using the same data.

This paper is organized as follows: First we will discuss the background of the project, including research on anomaly detection methods such as those used in NASA’s ACCEPT framework. We will then discuss the design of our system, beginning with our regression toolbox which implements a comparison of different regression algorithms, to our detection toolbox which implements a comparison of different detection methods. Next, we will present our results, demonstrating sufficient success in building an effective mechanism for detecting anomalous cold events. Finally, we will conclude and discuss future work for improving our system.

2. BACKGROUND (*DIANE/JENNA*)

Predicting adverse thermal events in this smart building lies in the model of anomaly detection since cold events occur as anomalies. There are a variety of methods which can be used to address the problem. First, we will discuss basic ideas relating to anomaly detection and then introduce NASA’s ACCEPT framework, which our system design is based off of. We will also discuss the related research from Cornell university which provides a good comparison for our work.

2.1 Anomaly Detection (*Diane*)

The general architecture of all anomaly detection methods consists of the following basic three modules or stages: pa-

*Team 2, Project C3, Host: Rodney Martin

¹<https://github.com/fedep3/sdl>

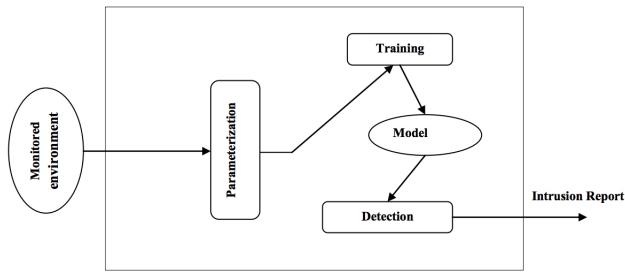


Figure 1: Generic Anomaly Detection Functional Architecture

parameterization, training and detection (see Figure 1).

Parameterization aims at collecting raw data from a monitored environment. The raw data should be representative of the system to be modeled, (e.g. Packet data from a network).

The training stage seeks to model the system using manual or automatic methods. The behaviors represented in the model will differ accordingly with the technique used.

Detection compares the system generated results in the training stage with the selected parameterized data portion. Threshold criteria will be selected to determine anomalous data events[16].

2.1.1 Anomaly Detection Techniques

There are three types of anomaly views. Supervised view: anomalies are what some user labels as anomalies. The data comprises of fully labeled training and test datasets. An ordinary classifier can be trained first and applied afterwards. The most common supervised algorithms are, Supervised Neural Networks, Support Vector Machines(SVM), k-Nearest Neighbors, Bayesian Networks and Decision Tree [12].

Unsupervised view: anomalies are outliers (points of low probability) in the data. The idea is evaluating the data solely based on intrinsic properties of the dataset. Furthermore, there is also no distinction between a training and a test dataset. Typically, distances or densities can be used to give an estimation what is normal and what is an outlier[17].

Semi-supervised view: in this situation, the training data only consists of normal data without any anomalies. The basic idea is that a model of the normal class is learned and anomalies can be detected afterwards by deviating from that model. Of course, in general any density estimation method can be used to model the probability density function of the normal classes, such as Gaussian Mixture Models[11] or Kernel Density Estimation[21].

Our research only focuses on semi-supervised anomaly detection setup. And ACCEPT, which our system is modeled after, is based on the Multivariate State Estimation Technique(MSET).

2.1.2 Multivariate State Estimation Technique

MSET[24] is a nonlinear, nonparametric modeling method that was originally developed by Argonne National Laboratory (ANL) for high-sensitivity proactive fault monitoring applications in commercial nuclear power applications. The MSET software described is configured to detect instrument degradation and perform system diagnostics with higher accuracy and faster response time than prior art techniques.

The implementation of this technique requires a few steps shown as follows:

Training step: A training procedure is developed to characterize the monitored equipment and build a model. There are some characteristics about Training Data: a) Containing all modes and ranges of operation b) Not containing any operating anomalies. After a comprehensive and error-free set of training data has been assembled, the MSET training algorithms are used to build an MSET model of the asset. The training procedure evaluates the training data and selects a subset of the training data observations (training vectors) that are determined to best characterize the asset's normal operation.

Monitoring step: This step uses the previously trained MSET model to estimate the expected values of the signals via newly acquired observations. A weighting method is used to produce the estimate by combining the example data values. Those examples most similar to the current observation are heavily weighted while those that are dissimilar are negligibly weighted. Similarity between the current observation and the learned examples is computed using sophisticated multivariable pattern matching techniques. The weighted combination of the most similar learned examples is used to compute the estimated signal values given the current observed signal values. After prediction, residuals can be obtained through computing the difference between a signal's predicted value and its directly sensed value. The MSET technique provides an extremely accurate estimate of sensor signals, with error rates that are typically 1% to 2% of the standard deviation of the input signal, which is excellent.

Fault detection step: This step uses Sequential Probability Ratio Test (SPRT) technique to determine whether the residual error value is uncharacteristic of the learned process model and thereby indicative of a sensor or equipment fault. For sudden, gross failures of a sensor or a subsystem, this procedure enunciates the disturbance as fast as a conventional threshold limit check. In operation, a time series of residual values are evaluated to determine whether the series of values is characteristic of the expected distribution or alternatively of some other specified distribution. Four possible fault-type distributions are considered in the current software. These are: 1) the residual mean value has shifted high; 2) the residual mean value has shifted low; 3) the residual variance value has increased; and 4) the residual variance value has decreased.

NASA's ACCEPT framework that our research based on is developed according to MSET model. They share some similarities and have differences.

2.2 ACCEPT Framework (Jenna)

ACCEPT stands for Adverse Condition and Critical Event Prediction Toolbox, and is an architectural framework developed by NASA for predicting anomalous events in real-time processes [18]. It is designed to utilize and compare a variety of machine learning techniques in order to provide a performance assessment of results for various applications. It is modeled after MSET (Multivariate State Estimation Technique), which is a state of the art statistical modeling technique originally developed by Argonne National Laboratory for online nuclear power plant monitoring [24]. MSET and ACCEPT share a similar fundamental basis: creating an accurate model of an asset from a sample of its normal operating data, and comparing the output from that model to the actual observation. Deviations detected from the model are used to predict process anomalies. MSET and ACCEPT differ, however, in that ACCEPT uses a variety of machine learning techniques to parameterize the underlying models rather than the more restrictive set used by MSET. Therefore, ACCEPT is a flexible system and an ideal model to observe for the problem we are exploring in this paper, where cold events are a rare occurrence.

2.2.1 Architecture

ACCEPT architecture (see Figure 2) is divided into Regression and Detection steps, with two distinct goals of finding the best-fitting regression model, then using the resulting predictions to configure the detection system and generate alarm statistics. All data is preprocessed and filtered using z-score normalization and feature selection. The training data, or the optimal data used that does not contain anomalous events, is provided to the regression toolbox to compare various static regression models to best fit the optimal system operation. Autoregressive models are not currently implemented in the architecture. All data, including the training data, validation data and testing data (only the first not containing anomalous events), is provided to the detection toolbox to configure the alarm system. The detection toolbox receives the residual prediction values given by

$$r_t = |y_{t,pred} - y_{t,actual}|$$

which represent the absolute difference between the regression algorithm output and the actual value of the output at a time t . These are generated over time and used to create various alarm statistics, with an Optimal Alarm System, Predictive Alarm System, SPRT Hypotheses and Redline Alarm System mechanisms implemented. The output from this system will be a binary classification indicating that an anomalous event will occur in a specific future time horizon or not.

2.2.2 Detection Models

In the detection toolbox, the residuals are used as the basis for learning the underlying linear dynamical system which implicitly models the dynamics that are unable to be captured by the static regression model. Capturing these unmodeled dynamics allows the system to test various statistical hypotheses used by the alarm systems. These unmodeled dynamics are indicated on the signal flow diagram (see Figure 3) via the dotted lines.

For the purposes of our project, we are focusing on studying

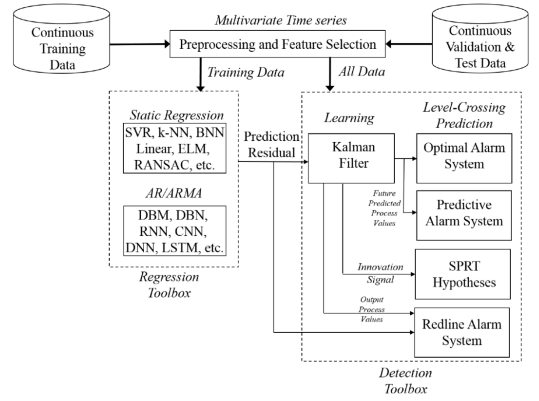


Figure 2: ACCEPT Functional Architecture

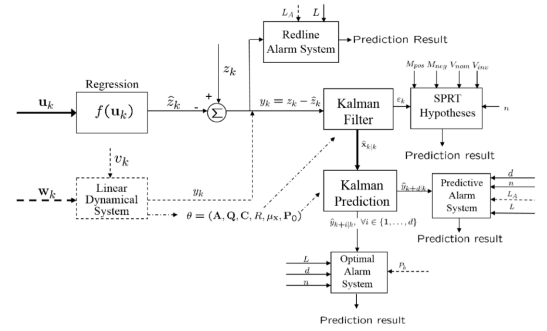


Figure 3: ACCEPT Signal Flow Diagram

the redline and predictive alarm systems, which use a level-crossing threshold mechanism to initiate an alarm. The difference between the two systems is that the redline alarm system simply uses the current prediction residual values directly as an early warning for prediction of an anomalous event. The predictive alarm system uses future predicted residual values generated from the use of a Kalman Filter in order to determine whether or not an anomalous event will occur in some time horizon. A Kalman Filter predicts future state estimations through Bayesian inference and an estimation of a joint probabilities in each time frame [13].

The formal definition for a level-crossing event is given by:

$$|y_{k+d|k}| > L_a$$

where y represents the prediction for d timesteps in advance given the current state k . From Figure 3, we can see that the predictive alarm system is supplied several configuration parameters, including the state dimension n and the future prediction horizon d . It is also supplied level-crossing thresholds L, L_a which are used to determine the residual value at which an alarm should be generated. These parameters may be chosen based on ROC Curve Analysis, which proceeds in a two-step process. First, the parameters dimension n and horizon d are selected based on maximizing the AUC of the ROC curve generated using the validation data. The second step is finding the ideal threshold(s) via the EER (Equal Error Rate) point, where threshold(s) are selected based on the point where the probability of a false alarm equals the

probability of a missed detection, denoted by:

$$P_{fa} = P_{md}$$

False alarms and missed detections are defined according to the following definition taken from the ACCEPT architecture documentation:

False Alarm: An alarm is triggered at a time point that does not contain an example of a confirmed anomalous event in at least one time point in the next d time steps.

Missed Detection: No alarm is triggered at a time point where an example of a confirmed anomalous event exists in at least one time point in the next d time steps.

The testing data is then used to generate statistics regarding the accuracy of the configured system. For more details on ACCEPT, please refer to [18].

2.3 Cornell University Research (*Victor*)

Being that we are basing our project off of the NASA ACCEPT framework, we are extending and comparing our results to the work done previously by Brutsaert et al. from Cornell University who used ACCEPT to predict adverse thermal events in NASAs Sustainability Base as well [4]. Because of this, we are attempting to follow in their footsteps with regards to the same values that they input as well as the same targets they optimized for. Our data is primarily the same data that they use, which is out of the possible 2544 sensor variables in the given dataset we are using the same 10 randomly selected features as they are to prevent any possible differences in our results due to training and testing our separate data. In their research, they also fixed a number of target variables such as the detection time as a fixed number of time steps in the future for the detection of adverse events (12 time steps or 1 hour) as well as the targeted temperature that would be considered an adverse event (below 68.1 degrees Fahrenheit). We adopt these constants as well as the dataset that they use.

Their results are also compelling to discuss as they found their best results to be from a Linear Regression model with a Predictive-Training Alarm System, Extreme Learning Machine model with a Predictive-Training Alarm System, and Extreme Learning Machine model with a Predictive-Validation Alarm System. Unlike ACCEPT, our system does not implement the use of training data in the detection phase, therefore all of our results are Validation based. We will go in-depth when we compare our results with theirs, but it is worth noting that surprisingly a model considered to be a relatively simple and basic machine learning method, Linear Regression, did very well in minimizing their False Alarm Rate and Missed Detection Rate.

3. REGRESSION TOOLBOX (*DIANE/VICTOR*)

We begin our discussion of system design by discussing the regression toolbox of our system. This includes preprocessing of the data and various regression models which are compared using cross-validation.

3.1 Preprocessing (*Diane*)

Our data collected from different sensors have different ranges. Some fluctuate around 5.0, some fluctuate around 1000.0. In

order to improve the accuracy of our model performance and find the precise coefficient between each feature, we tried to do some preprocessing that mainly focus on standardization.

3.1.1 Z-score Normalization

The result of standardization (or Z-score normalization) is that the features will be rescaled so that they'll have the properties of a standard normal distribution [25] with

$$\mu = 0, \sigma = 1$$

where μ is the mean (average) and σ is the standard deviation from the mean; standard scores (also called z scores) of the samples are calculated as follows:

$$z = \frac{x - \mu}{\sigma} \quad (1)$$

Standardizing the features so that they are centered around 0 with a standard deviation of 1 is not only important if we are comparing measurements that have different units, but it is also a general requirement for many machine learning algorithms. Intuitively, we can think of gradient descent as a prominent example (an optimization algorithm often used in logistic regression, SVMs, perceptrons, neural networks etc.). With features being on different scales, certain weights may update faster than others since the feature values x_j play a role in the weight updates

$$\Delta w_j = -\eta \frac{\partial J}{\partial w_j} = \eta \sum_i (t^{(i)} - o^{(i)}) x_j^{(i)} \quad (2)$$

3.1.2 Min-Max Scaling

An alternative approach to Z-score normalization (or standardization) is the so-called Min-Max scaling (often also simply called "normalization" - a common cause for ambiguities). In this approach, the data is scaled to a fixed range - usually 0 to 1. The cost of having this bounded range - in contrast to standardization - is that we will end up with smaller standard deviations, which can suppress the effect of outliers.

A Min-Max scaling is typically done via the following equation:

$$X_{norm} = \frac{X - X_{min}}{X_{max} - X_{min}} \quad (3)$$

In the ACCEPT model, it uses z-score scaling to gain better performance. In our research, we tried different standardization methods on different model to get the best performance separately by comparing their lowest median error like min-max standardization for Linear Regression model, z-score for SVM model and none at all in the case of RANSAC.

3.2 Regression Algorithm Comparison (*Victor*)

Just as in ACCEPT the goal of these various regressive models is to model the normal state of the system, thus it is only being run on the training data which does not contain any anomalous events. We partition the training data into several folds for the purpose of k-fold cross-validation analysis and based off of the results select a model minimizing the mean squared error.

In our case, we have considered 6 different regression methods. These being Linear Regression (LR), Extreme Learning Machines (ELM), and Random Sample Consensus (RANSAC), Support Vector Machines (SVM), k-Nearest Neighbors (KNN), and Bagging neural networks (BNN).

Linear Regression, Extreme Learning Machines, and Random Sample Consensus were all selected in order to compare our results to those in the Cornell paper as these three models were the models tested using ACCEPT. However, they each still have unique benefits with regard to possible data inputs.

For Support Vector Machines, k-Nearest Neighbors, and Bagging Neural Networks we discuss their benefits and the reason we selected each model in their respective sections.

3.2.1 Linear Regression

Linear regression is a simple forecasting model where through the use of a number of independent variables to fit a line to predict future values[4].

3.2.2 Extreme Learning Machine

ELMs have a structure similar to those of single layer feed-forward neural networks but the input layer parameters are assigned randomly. They are special in that training them takes approximately as long as a linear model but also have good generalization performance[4].

3.2.3 Random Sample Consensus

RANSAC is an iterative method that accounts for cases in which the assumptions of typical regression methods do not hold. It is a robust regression method against gross errors and extreme outliers in the data [4].

3.2.4 Support Vector Machine

Support Vector Machines are a set of supervised learning methods, among which include those for regression. They work ideally under circumstances with high dimensional spaces and is memory efficient in its use of a subset of training points in the decision function [26]. This was selected in part due to the dimensionality of our dataset as we trained based off of 11 different variables but can easily be extended to include more of the over 2000 other variables that Cornell did not use.

3.2.5 k-Nearest Neighbors

k-Nearest Neighbors is a local method that samples nearby cases and predicts a target based off of a similarity measure. This works best for data that is relatively non-volatile and based off of our initial survey of the temperature data this regressive model is applicable.

3.2.6 Bagging Neural Networks

Bagging or bootstrap aggregating is an ensemble meta-algorithm designed to improve the stability and accuracy of "unstable procedures" such as neural networks [3]. In practicality, this algorithm performs an internal generalization error estimate during its work which is similar to cross-validation, increasing computation time but also reduces variance and helps to avoid overfitting.

4. DETECTION TOOLBOX (*JENNA/FEDERICO*)

In the following sections we discuss the design of our detection toolbox, of which the goal is to create an optimal configuration of a time series modeling process and alarm system to detect future anomalous events in real time. First we will discuss the design of our system, including how residual values from the regression toolbox are used to generate alarms representing anomalous events at some time t . Next, we will discuss the various time-series modeling algorithms used to generate future predictions.

4.1 Design (*Jenna*)

In the previous section, we discussed the regression toolbox, which is used to find the best-fitting regression model to fit ideal training data, which contains no anomalous events. For our test dataset, this model will be used to predict future temperature values based on sensor data over time. The purpose of the detection toolbox, is to create a prediction model for the hidden dynamics of the system that the regression model cannot detect. By observing the temperature predictions over time, and measuring the difference between the actual values of the temperatures, we can generate a residual data signal that can be used to predict anomalous events (see Figure 4). As was previously discussed in the ACCEPT architecture [18], the residual data signal can be tested directly to create what is known as a redline alarm detection system, or past residuals can be used to predict future residual values to detect anomalous events. This second kind of system was defined as a predictive alarm system, and has shown to be significantly more effective than the simple redline alarm system for predicting future anomalous events in our data [4]. Our system is modeled after this second kind of system, but unlike ACCEPT we chose to use a variety of time-series modeling algorithms and configurations, to find the best one to suit the data, as opposed to the Kalman filtering system they use. This was done for a variety of reasons, namely that a pure Kalman filter [13] requires knowledge of the state transition model and other parameters, and for that reason can be difficult to configure. We chose to use advanced time-series modeling algorithms that are more general purpose instead, allowing us to test a variety of them and their configurations easily to find the right one for our data.

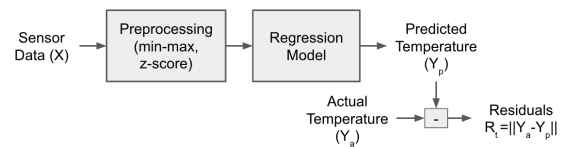


Figure 4: Regression Toolbox Signal Flow

4.1.1 Signal Flow

Figure 5 shows the signal flow diagram for the detection toolbox. At a given time t , past residual data of some length D is used to train the time series model to predict future residual values up to N steps in the future. These future residual values are provided to the predictive alarm system. The predictive alarm system uses the prediction at time $t+N$ in a level crossing mechanism, given by the equation below:

$$R_{t+N} > L_a$$

Where L_a is an arbitrary threshold value. In order to find the optimal parameters for D , N , L_a as well as the optimal time series model to use, the detection process proceeds as follows: First, various D , N and time series model combinations are tested, by generating an ROC curve using the spectrum of possible threshold values. The best models are selected based on their AUC (Area Under the Curve) for the ROC curve graph (see Figure 13). This is due to the fact that ROC AUC values are commonly used to predict the probability that a classifier will have a lower chance of error for a randomly chosen value from that curve, however they have been shown to be noisy metric [20]. The second step is taking the top models chosen and selecting a threshold for each based on EER (equal error rate). Errors are defined in a similar fashion to those of ACCEPT, wherein output classification is judged based on the ability to detect an anomalous event in the next N timesteps. Detection error tradeoff information is used to generate a DET (Detection Error Tradeoff) curve, and from that curve the threshold is selected that satisfied the criteria of $Pfa = Pmd$ (see Figure 14). The idea behind this step is to find the threshold that minimizes the total error of the system, therefore locating the point that minimizes both values at the same time. At this step, each of the models selected will be completely configured. Throughout this entire process, a set of validation data is used which contains anomalous events, unlike training data, to provide a more accurate fit to the real-world system. Finally, to test the system results are generated using the selected regression algorithm and configuration on the testing data to determine the accuracy of the system.

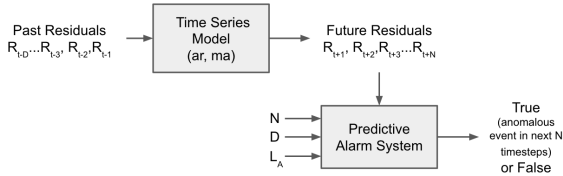


Figure 5: Detection Toolbox Signal Flow

4.1.2 Comparison to ACCEPT

As previously mentioned, our system was modeled after ACCEPT, but differs in the expansion to test various time-series generating algorithms and configurations instead of using a Kalman filtering mechanism. This includes the ability to configure a specific past prediction horizon instead of state space dimension. We also only have fully implemented a predictive alarm system mechanism, instead of a variety of alarm techniques, due to its promising accuracy in prior research [4]. We were also more coarse-grained in our search for correct configuration values, for instance we only used possible past prediction horizons of [16, 32, 48, 64, 80, 96] instead of the full spectrum of possible past horizons. We used methods like this which reduced the accuracy of our system, but allowed us to compute different model combinations in a reasonable amount of time (less than 24 hours) based on our empirical observations about the best data points to test.

4.2 Time-Series Modeling Algorithms *(Federico)*

The following are the various time series modeling algorithms used within our system.

4.2.1 ARIMA

Autoregressive moving average (ARMA) model is based on a combination of two processes, AR (autoregressive) and MA (moving average). The basic idea of an AR(p) model is that the current value, x_t , of a stationary series, can be explained as a linear function of p past values, $x_{t-1}, x_{t-2}, \dots, x_{t-p}$, where p determines the number of steps into the past needed to forecast the current value. The moving average model of order q , abbreviated as MA(q), assumes the white noise w_t up to lags q are combined linearly to form the observed data [14].

ARIMA resembles an ARMA model except that it is presumed that the time series has a steady underlying trend (see moving average). The model therefore works with the differences between the successive observed values, instead of the values themselves. To retrieve the original data from the differences requires a form of integration and the model is therefore called an autoregressive integrated moving average mode [29].

4.2.2 ARIMAX

In addition to past values of the response series and past errors, one can also model the response series using the current and past values of other series, called input series. The ARIMA model with input series, also called the ARIMAX model[10].

4.2.3 GARCH

Besides modeling the mean process using common ARIMA time series models, popularized by Box and Jenkins [2], it appears appropriate to incorporate non-linearities and fat-tailed distributions into volatility modeling.[31]

Autoregressive conditionally heteroscedastic (ARCH) models were introduced by Engle [8] and their GARCH (generalized ARCH) extension is due to Bollerslev [1]. In these models, the key concept is the conditional variance, that is, the variance conditional on the past. In the classical GARCH models, the conditional variance is expressed as a linear function of the squared past values of the series [9].

4.2.4 Gaussian State Space Models

State space models are used to represent time-varying systems, and consist of a latent state process $\{x_t\}_{t=1:T}$ and a related observation process $\{y_t\}_{t=1:T}$. The challenge is to infer the sequence of unknown latent states, and learn parameters of the model, from the sequence of known observations. It is usually assumed that the state process is Markovian (i.e., $x_t|x_{t-1}$ is independent of $x_{1:t-2}$) and that each observation depends only on the current state (i.e., $y_t|x_t$ is independent of $x_{1:t-1}$ and $x_{t+1:T}$) [5].

A basic linear state space model therefore obeys the following recursive system equations,

$$x_t = Fx_{t-1} + \epsilon_t^x \quad (4)$$

$$y_t = Hx_t + \epsilon_t^y \quad (5)$$

where $x_t \in \mathbb{R}^{d_x}$, $y \in \mathbb{R}^{d_y}$. $F \in \mathbb{R}^{d_x \times d_x}$ and $H \in \mathbb{R}^{d_y \times d_x}$ are the transition and observation matrices. $\epsilon_t^x \in \mathbb{R}^{d_x}$ and $\epsilon_t^y \in \mathbb{R}^{d_y}$ are the state disturbance and observation noise variables and we additionally assume that these are drawn from a Gaussian distribution,

$$\epsilon_t^x \sim \mathcal{N}(0, Q) \quad (6)$$

$$\epsilon_t^y \sim \mathcal{N}(0, R) \quad (7)$$

where Q and R are positive definite covariance matrices. This model can be written equivalently in terms of transition and observation densities,

$$p(x_t|x_{t1}, F, Q) = \mathcal{N}(x_t|Fx_{t1}, Q) \quad (8)$$

$$p(y_t|x_{t1}, H, R) = \mathcal{N}(y_t|Hx_{t1}, R) \quad (9)$$

The initial state x_1 may be known or may be assigned a Gaussian prior.

4.2.5 Vovk's Aggregating Algorithm

The Aggregating Algorithm (AA) [30] is a general approach to online learning that involves combining or 'merging' advice from a pool of experts (typically finite). The objective is to minimize the losses from a sequence of decisions that must be made in a stochastic environment. The convergence of the AA is moderated with a learning rate parameter that can be adjusted for each particular application but is otherwise constant. The Weak Aggregating Algorithm (WAA) is similar to the AA but uses a learning rate parameter that is proportional to \sqrt{n} [15].

5. RESULTS (FEDERICO/VICTOR/JENNA)

In this section we present the results of our experiments, using the Sustainable Building public data set [6]. This data set is composed of over 2544 different variables, of which we are utilizing data from 10 of them as predictors, and another sensor, "Building 232 Zone N240 room temperature," as our response variable [4]. The output variable we are observing is a temperature sensor representing the temperature of the primary room experiencing cold events (see Figures 6 and 7). Our goal in creating these results is to design a real time system capable of accurately predicting cold events, measured through the False Alarm (FA) and Missed Detection (MD) rate observed. To have a metric to determine the accuracy of our results vs. those of ACCEPT, we will be comparing our results to those of the Cornell Researchers who ran the same experiment using the same data [4].

Our complete system is composed of the regression and detection toolboxes, and we will be generating and comparing the results of each in order. Following similar logic as the Cornell University researchers, first an experiment to determine the best regression algorithm was run, including optimal preprocessing of the data. Then a second experiment

Index	Description	Object Type
A	232 RF1 HWS VALVE 14	Binary Output
B	232 A1 DX CAP SIGNAL	Analog Output
C	232 RSB P1 START/STOP	Binary Output
D	232 CRCP VALVE S28A	Analog Output
E	232 GWRV LOOPOUT	Analog Value
F	232 M1 AVG FLOW	Analog Value
G	232 ZONE N121 N125 AVERAGE TEM	Analog Value
H	232 S1 DPT AVG C	Analog Value
I	232 HP3 HEAT STAGE TIMER	Analog Value
J	232 N1 COOLING OFF	Analog Value

Figure 6: Input variables

Index	Description	Object Type
K	232 ZONE N240 ROOM TEMP	Analog Input

Figure 7: Output variable (Temperature)

was made to determine the best detection algorithm. Finally a complete run with the best combination of regression and detection algorithms was done.

5.1 Regression toolbox results (Victor)

As stated before, six different models were tested, Linear Regression (LR), Extreme Learning Machines (ELM), Random Sample Consensus (RANSAC), Support Vector Machines (SVM), k-Nearest Neighbors (KNN), and Bagging regressor (BNN). We first ran the data through cross-validation without any sort of pre-processing done to the data as shown in Figure 8. With this, we found similar results to those found by Cornell in that LR performed the best. It's worth noting that our diagram represents the negative mean squared error on the y axis as opposed to mean squared error, so lower error corresponds to higher positioning on the graph.

After this, we ran the data through various different types of pre-processing to see if the mean squared error could be decreased. We found the lowest median error with SVM, ELM, and LR when using Min-Max scaling, and for KNN and BNN we found the lowest with Z-score normalization. For RANSAC, it was none at all (this is due to the fact that RANSAC itself already performs some probabilistic preprocessing on the data within its run). Using these preprocessing pairings, we ran the algorithms through another 10-fold cross-validation on the training data. Comparing the results in Figure 9 to the original results in Figure 8 we can see that preprocessing does indeed slightly improve the results of each regression model, except in the case of SVM where preprocessed training data performed significantly better. SVM was ultimately the best algorithm in this final comparison.

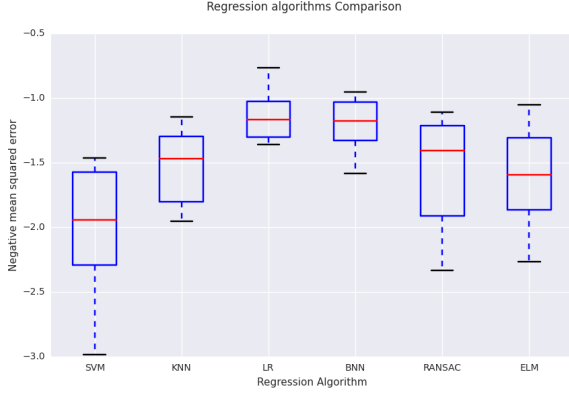


Figure 8: Regression Algorithm Comparison, before Preprocessing

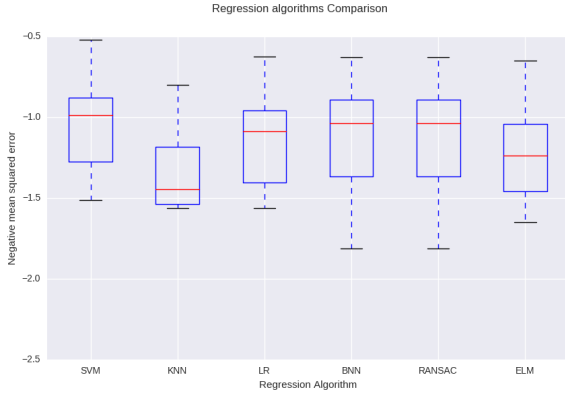


Figure 9: Regression Algorithm Comparison, after Preprocessing

5.2 Detection Toolbox Results *(Federico/Jenna)*

For choosing the best detection model, many parameters had to be configured in order to compare various time-series models and alarm parameters. As shown in Figure 5, the detection toolbox takes several parameters, including N , D , and L_a for configuring the predictive alarm system, and there are also many different time-series modeling algorithms to choose from. Most of the time-series modeling algorithms take additional parameters, denoted by p and q in the following results. In order to select the best algorithms, empirical observation was used to find a small collection of models and configurations. Since the performance of these models varied when used with various regression models, a variety were ultimately tested.

The following six different algorithm configurations were picked:

- ARIMA($p=1$, $q=1$)
- ARIMA($p=1$, $q=0$)
- ARIMA($p=2$, $q=0$)
- ARIMA($p=4$, $q=0$)

- GARCH($p=1$, $q=1$)
- Gaussian State Space Model (GGSM)

For future consideration of time constraints on a real system, each of these models was compared for time performance. The results of algorithm performance by type can be seen in Table 9 located in Appendix B.

Each algorithm was tested with a past prediction horizon in 32, 48, 64, 80, 96. These horizons were also determined empirically, through observation that most models performed most optimally around the middle of this range. The future prediction horizons that were tested include 6, 12, 18. Our purpose in testing several different future prediction horizons was to observe the relationship between this horizon and the accuracy of the observations. In a real-world system, this parameter may be fixed, depending on the usefulness of receiving an alarm about an impending event in different numbers of future timesteps. Since we are comparing our results to those of the Cornell researchers [4], we tested the same horizon they chose, of 12 future timesteps which represents one hour of actual time due to the sensor readings being from every 5 minutes. We then chose 6 future timesteps and 18 future timesteps for comparison to shorter or longer horizons, to see if this parameter would have a significant effect on accuracy. Finally, each one of these combinations of the aforementioned models and possible parameters was compared using the three best regression models from the regression toolbox results: Support Vector Machine, Linear Regression and Bagging Regressor.

Using these regression models, the residuals over time were computed for the validation data set, which represent the absolute difference between the output of the algorithm and the correct value.

With the predicted residuals, an ROC curve was generated and from it the area under the curve (AUC) value. At the same time, the ideal threshold was picked based on EER rate. Finally, the algorithms were sorted in by AUC in descending order, with a larger value corresponding to better performance.

The algorithmic representation of this portion of the detection toolbox may be seen in the following pseudocode:

Algorithm 1 Detection Toolbox Experiment

```

1: for algorithm  $\in$  top_algorithms do
2:   AUCvalues = []
3:   for model  $\in$  time_series_models do
4:     for past_horizon  $\in$  past_horizons do
5:       for future_horizon  $\in$  future_horizon do
6:         AUCvalues.append(computeROC-
          AUC(algorithm, model, past_horizon, future_horizon))
7:   print sorted(AUCvalues)

```

The full results from this experiment are shown in figures 10, 11 and 12 with the data used located in tables 3, 4 and 5 of Appendix B. The results for the top AUC values for each algorithm and prediction horizon have been graphed and grouped by algorithm. Each of the bar colors represents a different future prediction horizon. As we can see

from the graphs, lower future prediction horizon generally corresponds to greater AUC value. We can also see that the performance of the model varies depending on the regression algorithm used to create the residual data over time. All three algorithms overall showed similarities in accuracy when configured using the toolbox settings, but the optimal settings varied for each.

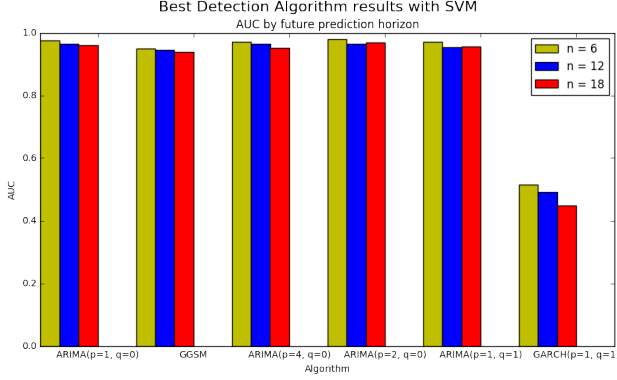


Figure 10: Best Detection Algorithm results with SVM

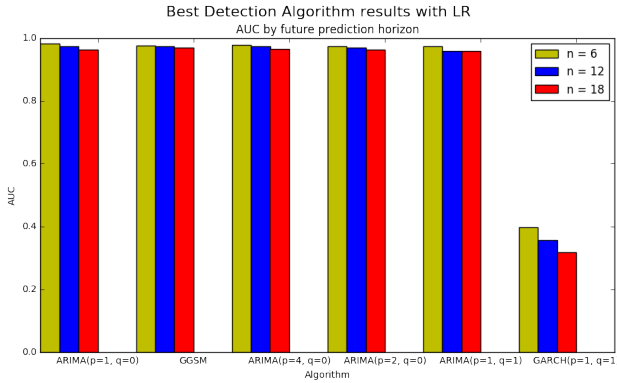


Figure 11: Best Detection Algorithm results with LR




Figure 12: Best Detection Algorithm results with BNN

5.2.1 ROC curve analysis

Conventionally, the performance of a classification rule is summarized by quantities related to the two types of errors involved in the decision process: true-positive rate and false-positive rate. The true-positive rate is the probability that a subject within the studied group is correctly classified within the group (the true-positive rate is also called sensitivity). The false-positive rate is the probability that a subject outside the group is incorrectly classified within the group (one minus the false-positive rate is called specificity).

The receiver operating characteristic is a plot of the true-positive rate (i.e., the ability of the test to detect the characteristic) versus the false-positive rate (i.e., the inability of the test to recognize a normal subject, without the studied characteristic, as normal) for all possible classification thresholds [19].

The area under the ROC curve (AUC) is often used in order to summarize the accuracy of a diagnostic system. In our system it was used to select the best detection algorithm. An example of an ROC curve generated in our analysis is shown in Figure 13, with the line of no discrimination shown. The line of no discrimination represents the ROC curve of a completely random classifier, used as the baseline comparison.

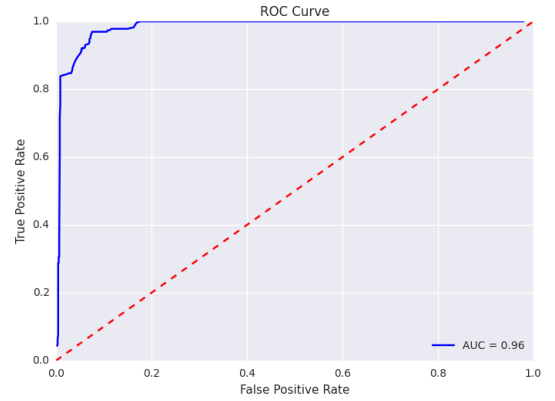


Figure 13: ROC (Receiver Operating Characteristic) for ARIMA(p=1,q=1), D=32, N=6

5.2.2 Detection Error Tradeoff

The false-negative rate or missed-detection rate (MDR), is the probability in a binary classifier that a object is classified as 0 (false) when the correct value was true. The false-positive rate (FPR), is the opposite case, where the object is classified as 1 (true) when the correct value was false. If we plot these values over different thresholds, the point, where both are the equal, is called the Equal Error Rate (EER). We can see in Figure 14 and example of a DET graph with the EER selection line given. We use this criteria to select the detection algorithm optimal threshold.

5.3 Complete Run Results (Federico/Jenna)

For each combination of regression and detection algorithm tried, the 3 best different algorithms for each future horizon were picked and run with the final testing data. The testing

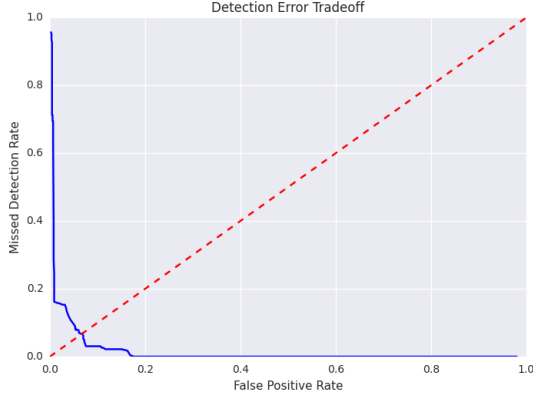


Figure 14: DET (Detection Error Tradeoff) for ARIMA($p=1,q=1$), $D=32$, $N=6$

data represents two continuous days from the same time frame, containing anomalous events. The results for this run are shown in figures 15, 16, 17, of which the data can be found in tables 6, 7, 8 in Appendix B. The performance of the models is shown with the False Alarm Rate (FAR) on the x-axis and the Missed Detection Rate (MDR) on the y-axis. The priority of optimizing for each of these metrics may vary depending on the given application, some systems may require very low MDR or very low FAR. For our application, we assumed that the goal was to minimize the sum of FAR and MDR, to reduce the overall error rate. We analyze our results accordingly, considering results ranked by this sum for performance.

5.4 Analysis (Federico/Jenna)

Looking at tables 6, 7 and 8, the best result overall was 0.6% FAR and 3.1% MDR, and was obtained using BNN, GGSM, past prediction horizon of 64, future prediction horizon of 6 and threshold of 3.225. Its worth noting there were also favorable results seen for other prediction horizons using the same model. We observed a 0.7% FAR and a 6.0% MDR using the same model but with a future prediction horizon of 12. For the future horizon of 18, we observed 0.7% FAR and 8.8% MDR. This shows overall strong performance of this model, unlike most of the other models which showed more variance in success at different time horizons.

5.4.1 Analysis of Time Series Models

In tables 3, 4 and 5, top results are inversely proportional to the future prediction horizon. This is also true with the testing data. GARCH was the worst model and in all tests we did it did not generate good results. In contrast all the ARIMA models, in particular the ones with $q = 0$, performed really well, same as GGSM. The last one is the most promising as it does not require parameter optimization. However, as indicated in Table 9, this model also takes the longest to generate, due computing demand required of computing unknown latent states.

Another discovery from these results was that the past prediction horizon needs to be equal or bigger than 48. In addition, all the threshold values discovered were between 2.5

and 4, which give us a rough estimate where the anomaly threshold will be for this dataset.




Figure 15: Best Run Results with SVM

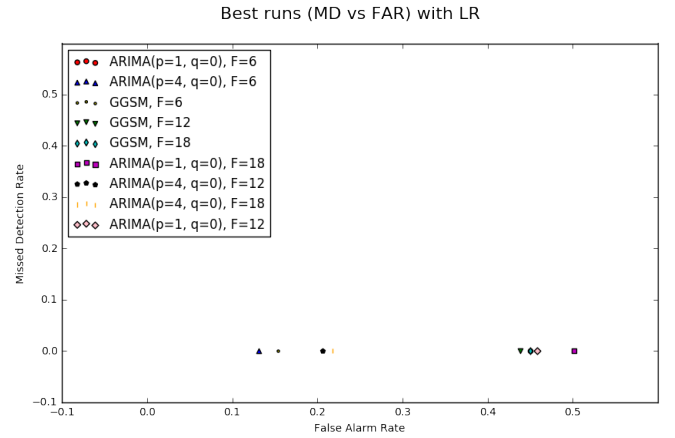


Figure 16: Best Run Results with LR




Figure 17: Best Run Results with BNN

Table 1: Best Results Overall for Future Horizon of 12

Method	False Alarm Rate	Missed Detection Rate
BNN w/ GGSM	0.007	0.060
BNN w/ ARIMA(p=4, q=0)	0.043	0.072
BNN w/ ARIMA(p=1, q=1)	0.043	0.078

Table 2: Cornell University Results

Method	False Alarm Rate	Missed Detection Rate
LR w/ Predictive-Training Alarm System	0.035	0.019
ELM w/ Predictive-Validation Alarm System	0.0138	0.0286
ELM w/ Predictive-Training Alarm System	0.0287	0.0333

5.4.2 Analysis of Regression Models

BNN overall had a very strong showing, with all of the top results being under about 10% for FAR and MDR rates. SVM gave the worst results, followed by LR. Further research needs to be done for SVM as it was the best with the validation data given, but failed to show significant accuracy with the testing data. The results with LR were not surprising due to the simplicity of the model. It is important to notice that SVM and LR both had missed detection rate 0 and high false alarm rate. This is caused because the threshold computed with the validation data is low, thus it does not miss any adverse event, but some of them are false.

5.4.3 Time Performance Analysis

When doing the experiments we also were able to determine the time efficiency each of the algorithm. As there are many factors that can have influence over the design of a system, we created Table 9 with a performance approximation that gives an idea on what timing to expect when doing a run with them on this data. As demonstrated by this table, the ARIMA model is the fastest to generate, with only about 3 minutes required to compute the necessary residual predictions on a 2.7 GHz Intel Core i5 machine. GGSM takes the longest, taking approximately 50 minutes to compute the same information. This information may be taken into account when designing a real system, where events may be more frequently generated and computing resources may be limited.

5.5 Comparison with Prior Research (Jenna)

Our results rival the ones obtained by Cornell University research team [4]. This may be seen in the summary of their top results in Table 2 and the summary of our top results with the same future prediction horizon in Table 1. We chose to compare based on the same prediction horizon, in order to accurately demonstrate the performance of our system under the same system performance constraint.

We can see that their best results ended up being ELM w/ a Predictive-Validation Alarm system, resulting in false alarm and missed detection rates of 1.38% and 2.86%. Our best results came from a BNN-GGSM model with a past prediction horizon of 64, resulting in false alarm and missed detection rates of 0.6% and 3.1% respectively. This demonstrates that our system may perform just as well as ACCEPT, for this particular data set. One thing that is curious, is the fact that ELM ended not performing very well in our regression toolbox tests, while ELM ended up being the best

performing regression model in the ACCEPT toolbox. This is likely caused by the underlying configuration of the model, in which the defaults were used and no additional configuration was performed to optimize the hyper-parameters of the model. This may have had a significant impact since the number of hidden nodes has a significant impact on ELM performance. Too many nodes can lead to overfitting while too few nodes may lead to underfitting [22]. Future work should include a method to configure these hyperparameters in the regression toolbox to further improve the performance of this model.

6. CONCLUSION (VICTOR)

This paper details and constructs an alternative anomaly detection system to that of NASA ACCEPT through the use of open source Python libraries. In modeling ACCEPT, we created a regression toolbox with the help of scikit-learn in order to perform cross-validation and to generate residuals. We then fed these residuals into our detection toolbox which utilized a variety of time series methods from pyflux in order to predict future results and signal alarms. The results that we obtained were not only on par with those generated by Cornell using ACCEPT, they also managed to surpass them in minimizing a number of valued metrics such as Missed Detection Rate and False Alarm Rate. Thus we were successful in our goal of designing a prediction system similar to ACCEPT that is both easier to use and more readily available.

7. FUTURE WORK (VICTOR)

There are multiple directions that our project could be taken in order to extend and to improve upon multiple aspects of our prediction toolbox.

7.1 Improving Runtime

Firstly, an issue that we ran into while working on this project was the runtime of our system. Being written in Python, it ran relatively slowly and could be improved through optimization techniques such as parallelization to multiple threads in order to speed up computation. Due to some of these limitations with regard to computing power, we had to limit the number of models and corresponding parameters we could test in a reasonable amount of time.

7.2 Improving Results

The experiments that we ran to find our ideal models were coarse grained and could be made to more finely search

through the different values that we varied to tune and improve our models. There were multiple hyperparameters that we did not fully explore. In our regression models, we did not tune any hyperparameters to improve the models to our training cross-validation. For Linear Regression, we used ordinary least squares where we could have used a different version in which we vary regularization coefficients [4]. For our neural network based algorithms such as ELM and BNN, there are a number of hyperparameters that can be tuned, not limited to varying the number of hidden neurons [4]. For ELM specifically, the Cornell results were much better than ours due to this hyperparameter tuning, so it is especially interesting to pursue pruning techniques for the number of hidden neurons such as the Successive Projections Algorithm [22]. Then for RANSAC, the cost threshold can be varied to determine the maximum deviation attributable to noise [4]. For kNN, the obvious hyperparameter that can be tuned is k . Then for SVMs, there are two primary parameters that can be varied to improve the model, those being the regularization parameter to minimize training error and model complexity and the Gaussian kernel that implicitly defines the nonlinear mapping from input to feature space [7].

7.3 Alternate Directions

Outside of improving our models, there are a number of other things that could be pursued within the scope of this project. There are different parameters that we could optimize for that we could fix while experimenting as to which time series model performs the best under those circumstances, the future prediction horizon specifically is a parameter that can be fixed. While we were attempting to minimize the Equal Error Rate, we could also specifically aim for a certain MDR or FAR while minimizing the other as an example in a real system it may be important to nearly never miss an alarm making MDR more valuable.

7.4 Other Projects

For others interested in using our prediction system, all code is open source and publicly available. Because of this, multiple changes to the code and system are possible. Rather than specifically being applied to temperature and cold events, the code could be changed to accommodate other applications. Further documentation would be important for these goals. See Appendix A for information on accessing, installing and running our software.

8. REFERENCES

- [1] T. Bollerslev. Generalized autoregressive conditional heteroskedasticity. *Journal of Econometrics*, 31(3):307–327, 1986.
- [2] J. G. Box GEP. *Time series analysis: forecasting and Control*. Holden-Day, San Francisco, 1970.
- [3] L. Breiman and L. Breiman. Bagging predictors. In *Machine Learning*, pages 123–140, 1996.
- [4] H. Brutsaert, Y. Gao, Y. Rim, and Y. Xie. Using nasa accept (adverse condition and critical event prediction toolbox) for adverse thermal events in a smart building. 2016.
- [5] P. Bunch, J. Murphy, and S. Godsill. Bayesian learning of degenerate linear gaussian state space models using markov chain monte carlo. *IEEE Transactions on Signal Processing*, 64(16):4100–4112, Aug 2016.
- [6] Data from sustainability base characterizing cold complaints for accept, 2016. Available at c3.nasa.gov/dashlink/resources/936/.
- [7] K. Duan, S. S. Keerthi, and A. N. Poo. Evaluation of simple performance measures for tuning svm hyperparameters. *Neurocomputing*, 51:41–59, 2003.
- [8] R. Engle. Autoregressive conditional heteroscedasticity with estimates of the variance of united kingdom inflation. *Econometrica*, 50(4):987–1007, 1982.
- [9] C. Francq and J.-M. ZakoÁřan. *GARCH(p, q) Processes*, pages 17–61. John Wiley & Sons, Ltd, 2010.
- [10] Y. Hao, J. Wu, Q. Sun, Y. Zhu, Y. Liu, Z. Li, and T.-C. J. Yeh. Simulating effect of anthropogenic activities and climate variation on liulin springs discharge depletion by using the arimax model. *Hydrological Processes*, 27(18):2605–2613, 2013.
- [11] N.-C. R. C. S. in Probability and S. P. S. U. I. of Mathematical Statistics. Lindsay b. mixture models: Theory, geometry, and applications. 1995.
- [12] Q. JR. Programs for machine learning. 2002.
- [13] R. E. Kalman. A new approach to linear filtering and prediction problems. *Transactions of the ASME-Journal of Basic Engineering*, 82(Series D):35–45, 1960.
- [14] U. Kumar and V. K. Jain. Arima forecasting of ambient air pollutants (o3, no, no2 and co). *Stochastic Environmental Research and Risk Assessment*, 24(5):751–760, 07 2010. Copyright - Springer-Verlag 2010; Āłltima actualizaciĀłn - 2014-08-22.
- [15] T. Levina, Y. Levin, J. McGill, M. Nediak, and V. Vovk. Weak aggregating algorithm for the distribution-free perishable inventory problem. *Operations Research Letters*, 38(6):516 – 521, 2010.
- [16] Machine learning techniques for anomaly detection: An overview. *International Journal of Computer Applications*, 2016.
- [17] S. U. Markus Goldstein. A comparative evaluation of unsupervised anomaly detection algorithms for multivariate data. 2016.
- [18] R. A. Martin, S. Das, and S. Hosein. Accept: Introduction of the adverse condition and critical event prediction toolbox. 2015.
- [19] P. Martinez-Camblor, C. Carleos, and N. Corral. General nonparametric {ROC} curve comparison. *Journal of the Korean Statistical Society*, 42(1):71–81, 2013.
- [20] R. R. Maxion, R. A. Roberts. Proper use of roc curves in intrusion/anomaly detection. *Technical Report Series*, CS-TR-871, 2004.
- [21] T. A. of Mathematical Statistics. Rosenblatt m. remarks on some nonparametric estimates of a density function. 1956.
- [22] D. Parente Mesquita, J. Gomes, L. Ramos Rodrigues, and R. Kawakami Galvao. Pruning extreme learning machines using the successive projections algorithm. *IEEE Latin America Transactions*, 13(12):3974–3979, 2015.
- [23] Pyflux, 2016. Available at www.pyflux.com.
- [24] R. D. R. S. Randall B, Die Davis. 1 development of an

online predictive monitoring system for power generating plants. 2014.

- [25] S. Raschka. About feature scaling and normalization and the effect of standardization for machine learning algorithms. 2014.
- [26] Support vector machines. Available at www.scikit-learn.org.
- [27] Scikit-learn.org, 2016. Available at www.scikit-learn.org.
- [28] Sustainability base: A living testbed for nasa technologies, 2012. Available at www.nasa.gov/ames/facilities/sustainabilitybase.
- [29] G. Upton and I. Cook. *A Dictionary of Statistics (3 ed.)*. Oxford University Press, 2014.
- [30] V. Vovk. Competitive on-line statistics. *International Statistical Review*, 69(2):213–248, 2001.
- [31] A. C. M. Zeileberger and A. Brauneis. Modeling carbon spot and futures price returns with garch and markov switching garch models. *Central European Journal of Operations Research*, 24(1):149–176, 2016.

APPENDIX

A. SOFTWARE GUIDE (JENNA)

In the following section, we describe how to run our project software, hosted online on github at the following link:

<https://github.com/fedep3/sdl>

A.1 Software Installation

Since this project requires the installation of pyflux, matplotlib, scikit-learn and all of their dependencies, we recommend creating a virtual environment on your personal machine and installing all packages there in order to facilitate easy installation. If you are using a Mac OS, you must upgrade your operating system to the latest 2016 version, macOS Sierra, due to Python library dependencies that conflict on older versions of the operating system. Our system requires the installation of Python 2.7, and the following instructions assume it has been installed and is the default version of python installed on your system.

The following instructions will create a virtual environment and install all dependencies required for our project software in a few easy steps:

Run the following command to install the virtual environment package:

```
sudo pip install virtualenvwrapper
```

Add the following lines to your .bashrc file:

```
export WORKON_HOME=~/.virtualenvs
export PROJECT_HOME=/path/to/sdl
source /usr/local/bin/virtualenvwrapper.sh
```

To create a new virtual environment: If only Python 2.7 is installed run: `mkvirtualenv sdl` Otherwise run:

```
mkvirtualenv --python=/path/to/python/2.7 sdl
```

where `/path/to/python/2.7` represents the path to the python 2.7 executable on your machine

Finally, in order to facilitate easy installation we have listed out package requirements in the requirements.txt file, located in the github. Run the following command to install all packages within the file:

```
pip install -r requirements.txt
```

Now, each time a terminal is opened and you wish to work on the project, you must run the following command to transition to your new virtual environment:

```
workon sdl
```

If one wishes to not use a virtual environment and install the packages globally on a machine, simply run the previous pip command.

A.2 Running the Software

Note: Please refer to the online github link provided to find the most updated version of these instructions.

The software at the time of this writing has 3 modes: regression toolbox mode, detection toolbox mode, and best run mode. The `-t` command line option is used to specify what mode to run, and the optional argument `-r` is used to specify what regression algorithm to use. Use the `-O` python option to suppress debugging output and graph plotting, which is necessary when many different configurations are being run:

```
python [-O] main.py [-t type] [-r reg_alg]
```

A.2.1 Regression Toolbox Mode

To run the regression algorithm comparison and generate the corresponding box plot, use the `-t reg` option, which will run the data preprocessing and regression algorithm comparison and display a box plot representing the statistics on the negative mean squared error for observation.

A.2.2 Detection Toolbox Mode

The detection toolbox mode is designed to search for the ideal parameters for a given regression algorithm, and output the best findings ranked by the AUC value. To run this mode, use the `-t det` option, as well as the `-r` option used to specify which regression algorithm to use in the detection toolbox. For instance, to run the detection toolbox using SVM, use the `-r SVM` option.

A.2.3 Best Run Mode

This mode is designed to generate final results using testing data. To run it, use the `-t best` option. The best runs are currently hardcoded into this function, including the top results from the top three regression algorithms. The results are ranked based on lowest combined false alarm and missed detection rate.

B. FULL DATA TABLES

The following page contains the complete result data tables used to generate the graphs contained within this report.

Table 3: Detection Toolbox Top Results with SVM

Detection Algorithm	ROC AUC	Past P. H.	Future P. H.	Threshold	False Alarm Rate	Missed Detection Rate
ARIMA(p=2, q=0)	0.9786	64	6	2.575	0.033	0.044
ARIMA(p=1, q=0)	0.9752	48	6	2.925	0.048	0.048
ARIMA(p=1, q=1)	0.9718	64	6	2.575	0.041	0.044
ARIMA(p=2, q=0)	0.9690	80	18	2.575	0.076	0.089
ARIMA(p=2, q=0)	0.9651	80	12	2.525	0.082	0.078
ARIMA(p=4, q=0)	0.9650	80	12	2.575	0.064	0.070
ARIMA(p=1, q=0)	0.9642	64	12	2.525	0.075	0.073
ARIMA(p=1, q=0)	0.9601	80	18	2.575	0.074	0.077
ARIMA(p=1, q=1)	0.9569	80	18	2.575	0.083	0.085

Table 4: Detection Toolbox Top Results with LR

Detection Algorithm	ROC AUC	Past P. H.	Future P. H.	Threshold	False Alarm Rate	Missed Detection Rate
ARIMA(p=1, q=0)	0.9808	32	6	3.025	0.067	0.066
ARIMA(p=4, q=0)	0.9770	48	6	3.475	0.050	0.050
GGSM	0.9763	64	6	3.275	0.052	0.042
ARIMA(p=1, q=0)	0.9742	48	12	3.025	0.069	0.071
GGSM	0.9729	64	12	3.025	0.057	0.060
ARIMA(p=4, q=0)	0.9724	48	12	3.075	0.070	0.071
GGSM	0.9698	64	18	2.625	0.053	0.071
ARIMA(p=4, q=0)	0.9646	64	18	3.125	0.072	0.067
ARIMA(p=1, q=0)	0.9622	48	18	2.975	0.081	0.078

Table 5: Detection Toolbox Top Results with BNN

Detection Algorithm	ROC AUC	Past P. H.	Future P. H.	Threshold	False Alarm Rate	Missed Detection Rate
ARIMA(p=4, q=0)	0.9770	96	12	3.475	0.059	0.056
ARIMA(p=1, q=1)	0.9748	96	12	3.275	0.069	0.068
GGSM	0.9710	64	6	3.225	0.064	0.056
GGSM	0.9701	64	12	3.075	0.045	0.021
GGSM	0.9700	64	18	2.925	0.029	0.028
ARIMA(p=2, q=0)	0.9694	96	6	3.275	0.077	0.069
ARIMA(p=4, q=0)	0.9689	96	18	3.325	0.056	0.052
ARIMA(p=4, q=0)	0.9673	96	6	3.425	0.065	0.074
ARIMA(p=1, q=1)	0.9621	96	18	3.225	0.066	0.063

Table 6: Best Run Results with SVM

Detection Algorithm	Past P. H.	Future P. H.	Threshold	False Alarm Rate	Missed Detection Rate
ARIMA(p=1, q=0)	48	6	2.925	0.447	0.000
ARIMA(p=2, q=0)	64	6	2.575	0.521	0.000
ARIMA(p=1, q=1)	64	6	2.575	0.534	0.000
ARIMA(p=4, q=0)	80	12	2.575	0.538	0.000
ARIMA(p=2, q=0)	80	12	2.525	0.555	0.000
ARIMA(p=1, q=1)	80	18	2.575	0.557	0.000
ARIMA(p=2, q=0)	80	18	2.575	0.561	0.000
ARIMA(p=1, q=0)	80	18	2.575	0.564	0.000
ARIMA(p=1, q=0)	64	12	2.525	0.575	0.000

Table 7: Best Run Results with LR

Detection Algorithm	Past P. H.	Future P. H.	Threshold	False Alarm Rate	Missed Detection Rate
ARIMA(p=4, q=0)	48	6	3.475	0.131	0.000
GGSM	64	6	3.275	0.153	0.000
ARIMA(p=4, q=0)	48	12	3.075	0.206	0.000
ARIMA(p=4, q=0)	64	18	3.125	0.218	0.000
GGSM	64	12	3.025	0.439	0.000
ARIMA(p=1, q=0)	32	6	3.025	0.450	0.000
GGSM	64	18	2.625	0.450	0.000
ARIMA(p=1, q=0)	48	12	3.025	0.458	0.000
ARIMA(p=1, q=0)	48	18	2.975	0.502	0.000

Table 8: Best Run Results with BNN

Detection Algorithm	Past P. H.	Future P. H.	Threshold	False Alarm Rate	Missed Detection Rate
GGSM	64	6	3.225	0.006	0.031
ARIMA(p=4, q=0)	96	6	3.425	0.013	0.037
ARIMA(p=2, q=0)	96	6	3.275	0.022	0.037
GGSM	64	12	3.075	0.007	0.060
GGSM	64	18	2.925	0.007	0.088
ARIMA(p=4, q=0)	96	12	3.475	0.043	0.072
ARIMA(p=1, q=1)	96	12	3.275	0.043	0.078
ARIMA(p=4, q=0)	96	18	3.325	0.073	0.104
ARIMA(p=1, q=1)	96	18	3.225	0.073	0.110

Table 9: Algorithm Time Performance for 2.7 GHz Intel Core i5 Machine

Algorithm	Time Performance
ARIMA	3 mins
GARCH	30 mins
GSSM	50 mins
Aggregating	> 120 mins